

Splunk Core Certified User

Offline exam-cram workbook • 8 short sessions • beginner → exam-ready • aligned to the current official blueprint

YOUR LAB DEFAULT The examples use index=test because you already confirmed it works. If a field is absent, open one event, inspect its fields, and substitute a field that exists in your data.

How to study

Use one or two sessions per day. For each 25–35 minute session: learn (5 min) → run every search (15 min) → explain the output aloud (5 min) → answer the Quick Check without looking (5 min). Re-run searches with All time only when necessary; narrow time ranges are faster.

Exam map — put most time where the points are

22% Basic Searching	20% Using Fields	15% Search Language	15% Transforming Commands
12% Reports & Dashboards	6% Lookups	5% Splunk Basics	5% Scheduled Reports & Alerts

SESSION 1 • WHAT SPLUNK DOES + WHERE THINGS LIVE

Goal: explain the data path and navigate Search & Reporting. A forwarder sends machine data. An indexer receives, processes, indexes, and stores it. A search head lets users run searches and build reports/dashboards. An app is a packaged workspace containing searches, dashboards, configurations, and other knowledge for a use case.

RUN index=test | head 20

- Locate the Search bar, time picker, Events/Statistics/Visualization tabs, fields sidebar, timeline, and Job controls.
- Open one event and identify _time, _raw, host, source, sourcetype, and index.

QUICK CHECK Which component normally indexes and stores incoming data? **Answer: The indexer.**

SESSION 2 • BASIC SEARCHING + TIME

Goal: reduce a large event set precisely. Terms are implicitly ANDed. Use uppercase AND, OR, and NOT; quotes preserve an exact phrase; * is a wildcard. Prefer index and sourcetype filters early because they reduce the work Splunk must do.

```
RUN index=test error
index=test ("failed" OR "denied") NOT success
index=test earliest=-24h latest=now
index=test host=* sourcetype=*
```

Time: -15m = last 15 minutes, -24h = last 24 hours, @d = start of day, now = current time. Use the time picker or earliest/latest modifiers. Confirm the timeline changes when you change the time range.

QUICK CHECK Which is more efficient: status=failure, or index=test sourcetype=known status=failure? **Answer: The second; it constrains index and sourcetype.**

SESSION 3 • EVENTS + FIELDS

Goal: use fields as searchable name=value pairs. _raw is the original event text; _time is the event time. host identifies the originating system, source identifies the input/file/path, and sourcetype describes the data format. Selected fields are pinned by you; interesting fields are discovered because they appear often enough in results.

```
RUN index=test | table _time host source sourcetype
index=test | fields _time host source sourcetype
index=test host=* | fields - _raw
```

table keeps only the named fields and formats them as columns. fields includes (+) or excludes (-) fields without creating a statistical table. Click a field value in the sidebar to add it to the search.

QUICK CHECK Which metadata field usually identifies the event data format? **Answer: sourcetype.**

Build SPL skill

Sessions 4–8 turn raw events into defensible answers, reports, dashboards, enriched results, and alerts.

SESSION 4 • THE PIPELINE + RESULT COMMANDS

A pipe (|) sends the current results to the next command. Order matters: filter early, then shape the remaining rows. `sort - _time` means newest first; `dedup user` then keeps the first event for each user—so sorting before `dedup` controls which event survives.

```
SPL LADDER index=test
| fields _time host user status
| sort - _time
| dedup user
| rename host AS system
| table _time system user status
```

- table: display chosen columns • rename: change a field label • fields: keep/remove fields
- dedup: remove repeated field values • sort: order results (+ ascending, - descending) • head: keep first N rows

QUICK CHECK To keep the newest event per user, what should come before `dedup user`? **Answer: `sort - _time`.**

SESSION 5 • TOP, RARE + STATS

These are transforming commands: they convert events into a table of calculated results. `top` returns most common values (plus count and percent); `rare` returns least common values; `stats` calculates values such as count, distinct count, sum, average, minimum, and maximum. `BY` creates one row per group.

```
RUN index=test | top limit=5 host
index=test | rare limit=5 sourcetype
index=test | stats count BY host
index=test | stats dc(user) AS unique_users BY host
index=test | stats avg(bytes) AS avg_bytes max(bytes) AS max_bytes BY host
```

Function memory: `count` counts events; `dc(field)` counts distinct values; `sum/avg/min/max` need a numeric field. If `bytes` does not exist, choose a numeric field from your sidebar.

QUICK CHECK Which command answers ‘How many events occurred for each host’? **Answer: `stats count BY host`.**

SESSION 6 • REPORTS, VISUALIZATIONS + DASHBOARDS

A report is a saved search that can be run again and scheduled. A visualization is a graphical view of transformed results. A dashboard is a collection of panels, often powered by saved reports. Build from a working search: `Save As` → Report; choose a visualization; then Add to Dashboard and edit the panel/title/permissions as needed.

```
PRACTICE index=test | timechart span=1h count
index=test | stats count BY host | sort - count
```

Choose a table for exact values, bar/column for category comparisons, line/timechart for trends, and pie only for a few parts of a whole. A chart needs transformed, structured results—not just raw events.

QUICK CHECK What holds multiple panels in one view? **Answer: A dashboard.**

SESSION 7 • LOOKUPS

A lookup enriches event data with fields from an external table (commonly CSV). Example file: `host, owner, criticality`. Upload the file, create a lookup definition, optionally configure an automatic lookup, then use the definition name in SPL. `OUTPUT` adds matching columns; `OUTPUTNEW` avoids overwriting an existing field.

```
RUN index=test
| lookup asset_lookup host OUTPUT owner criticality
| table _time host owner criticality
```

QUICK CHECK What is the core purpose of a lookup? **Answer: Add context to events by matching a field to external data.**

SESSION 8 • SCHEDULED REPORTS + ALERTS

A scheduled report runs a saved search on a timetable and produces/updates results. An alert runs a search and triggers an action when a condition is met. Typical settings: schedule, time range, trigger condition, throttling, permissions, and action. `Fired Alerts` shows triggered alerts.

```
ALERT SEARCH index=test ("failed" OR "denied") | stats count AS failures
```

Practice: save the search as a report, schedule it hourly, then create a separate alert that triggers when `failures > 10`. Think carefully about the search time range versus the schedule—an hourly alert often examines the previous hour.

QUICK CHECK You need notification only when failures exceed 10. Report or alert? **Answer: Alert.**

Final review

Do the quiz closed-book, check the explanations, then use the readiness checklist to target your final study time.

10-question exam-style check

- | | |
|--|---|
| 1. Which component receives, indexes, and stores incoming data? A Search head
B Indexer C Dashboard D App | 6. Which SPL counts events for each status? A top count status B stats count BY
status C table status count D fields status |
| 2. Which syntax searches an exact phrase? A failed login B [failed login] C "failed
login" D failed*login | 7. Which function counts unique users? A count(user) B sum(user) C dc(user) D
rare(user) |
| 3. Which metadata field describes the data format? A source B host C index D
sourcetype | 8. You need the newest event per user. Which sequence is correct? A dedup user
 sort - _time B sort - _time dedup user C rare user head 1 D fields user top 1 |
| 4. Which command removes _raw from the returned fields? A fields - _raw B table
- _raw C dedup _raw D rare _raw | 9. What does a lookup primarily do? A Delete duplicates B Add external context C
Schedule a search D Build an index |
| 5. Which search is best constrained? A failure B status=failure C index=test
failure D index=test sourcetype=known status=failure | 10. Which object triggers an action when a search condition is met? A Report B
Dashboard C Alert D App |

Answer key + why

1 B—indexers index/store data. 2 C—quotes preserve a phrase. 3 D—sourcetype classifies format. 4 A—fields - excludes a field. 5 D—index + sourcetype + field filter narrows early. 6 B—stats aggregates and BY groups. 7 C—dc means distinct count. 8 B—sort first so dedup keeps the newest. 9 B—lookups enrich matching events. 10 C—alerts trigger actions.

SCORE GUIDE 9–10: ready for a full timed mock. 7–8: re-run the missed session. 0–6: repeat Sessions 2–5 before moving on.

SPL pocket reference

Find	index=test error • field=value • "exact phrase" • A OR B • NOT value
Time	earliest=-15m latest=now • earliest=-24h • earliest=@d
Shape	fields + a,b • fields - _raw • table a b • rename a AS label
Clean	sort - _time • dedup user • head 20
Summarize	top limit=5 host • rare user • stats count BY status • dc(user)
Trend/enrich	timechart span=1h count • lookup asset_lookup host OUTPUT owner

Final readiness checklist

- | | |
|--|---|
| ☐ I can explain forwarder → indexer → search head and define a Splunk app. | ☐ I can save/edit a report, choose a sensible visualization, and add/edit a dashboard panel. |
| ☐ I can narrow by time, index, sourcetype, keywords, fields, Boolean operators, quotes, and wildcards. | ☐ I can describe CSV upload → lookup definition → optional automatic lookup → lookup SPL. |
| ☐ I can identify _time, _raw, host, source, sourcetype, index, selected fields, and interesting fields. | ☐ I can distinguish a scheduled report from an alert and explain Fired Alerts and throttling. |
| ☐ I can predict the result of table, fields, rename, dedup, sort, top, rare, and stats. | ☐ I can complete 60 questions in 57 minutes (about 57 seconds each) and flag hard items for review. |
| ☐ I can write stats count BY field and dc(field), and explain why pipeline order matters. | ☐ I reviewed every missed question by explaining why each wrong option is wrong —not just memorizing the answer. |

Official reference points (verified July 2026)

Blueprint: splunk.com/en_us/pdfs/training/splunk-test-blueprint-user.pdf • Certification page: splunk.com/en_us/training/certification-track/splunk-core-certified-user.html • SPL command reference: help.splunk.com